

Contents

1 Basic Test Results	2
2 README	3
3 Barrier.h	4
4 Barrier.cpp	5
5 Makefile	6
6 MapReduceFramework.cpp	7

1 Basic Test Results

```
1  ===== Tar Content Test =====
2  found README
3  found Makefile
4  tar content test PASSED!
5
6  ===== logins =====
7  login names mentioned in file:  tomka,alonemanuel
8  Please make sure that these are the correct login names.
9
10 ===== make Command Test =====
11 g++ -Wall -std=c++11 -g -pthread -I.   -c -o MapReduceFramework.o MapReduceFramework.cpp
12 g++ -Wall -std=c++11 -g -pthread -I.   -c -o Barrier.o Barrier.cpp
13 ar rv libMapReduceFramework.a MapReduceFramework.o Barrier.o
14 a - MapReduceFramework.o
15 a - Barrier.o
16 ranlib libMapReduceFramework.a
17
18 ar: creating libMapReduceFramework.a
19
20 make command test PASSED!
21
22 Pre-submission passed!
23 Keep in mind that this script tests only basic elements of your code.
```

2 README

```
1 tomka, alonemanuel
2 Tom Kalir (316426485), Alon Emanuel (205894058)
3 EX: 3
4
5 FILES:
6
7 MapReduceFramework.cpp
8 Barrier.h
9 Barrier.cpp
10 Makefile
11
12 REMARKS:
```

3 Barrier.h

```
1  #ifndef BARRIER_H
2  #define BARRIER_H
3  #include <pthread.h>
4
5  // a multiple use barrier
6
7  class Barrier {
8  public:
9      Barrier(int numThreads);
10     ~Barrier();
11     void barrier();
12
13 private:
14     pthread_mutex_t mutex;
15     pthread_cond_t cv;
16     int count;
17     int numThreads;
18 };
19
20 #endif //BARRIER_H
```

4 Barrier.cpp

```
1  #include "Barrier.h"
2  #include <cstdlib>
3  #include <stdio>
4
5  Barrier::Barrier(int numThreads)
6      : mutex(PTHREAD_MUTEX_INITIALIZER)
7      , cv(PTHREAD_COND_INITIALIZER)
8      , count(0)
9      , numThreads(numThreads)
10 { }
11
12
13 Barrier::~Barrier()
14 {
15     if (pthread_mutex_destroy(&mutex) != 0) {
16         fprintf(stderr, "[Barrier] error on pthread_mutex_destroy");
17         exit(1);
18     }
19     if (pthread_cond_destroy(&cv) != 0){
20         fprintf(stderr, "[Barrier] error on pthread_cond_destroy");
21         exit(1);
22     }
23 }
24
25
26 void Barrier::barrier()
27 {
28     if (pthread_mutex_lock(&mutex) != 0){
29         fprintf(stderr, "[Barrier] error on pthread_mutex_lock");
30         exit(1);
31     }
32     if (++count < numThreads) {
33         if (pthread_cond_wait(&cv, &mutex) != 0){
34             fprintf(stderr, "[Barrier] error on pthread_cond_wait");
35             exit(1);
36         }
37     } else {
38         count = 0;
39         if (pthread_cond_broadcast(&cv) != 0) {
40             fprintf(stderr, "[Barrier] error on pthread_cond_broadcast");
41             exit(1);
42         }
43     }
44     if (pthread_mutex_unlock(&mutex) != 0) {
45         fprintf(stderr, "[Barrier] error on pthread_mutex_unlock");
46         exit(1);
47     }
48 }
```

5 Makefile

```
1  CC=g++
2  CXX=g++
3  RANLIB=ranlib
4
5  LIBSRC=MapReduceFramework.cpp Barrier.cpp
6  LIBOBJ=MapReduceFramework.o Barrier.o
7
8  INCS=-I.
9  CFLAGS = -Wall -std=c++11 -g -pthread $(INCS)
10 CXXFLAGS = -Wall -std=c++11 -g -pthread $(INCS)
11
12 LIBMAPREDUCE = libMapReduceFramework.a
13 TARGETS = $(LIBMAPREDUCE)
14
15 TAR=tar
16 TARFLAGS=-cvf
17 TARNAME=ex3.tar
18 TARSRC=$(LIBSRC) Makefile README Barrier.h
19
20 all: $(TARGETS)
21
22 $(TARGETS): $(LIBOBJ)
23     $(AR) $(ARFLAGS) $@ $^
24     $(RANLIB) $@
25
26 clean:
27     $(RM) $(TARGETS) $(LIBTHREADS) $(OBJ) $(LIBOBJ) *~ *core
28
29 depend:
30     makedepend -- $(CFLAGS) -- $(SRC) $(LIBSRC)
31
32 tar:
33     $(TAR) $(TARFLAGS) $(TARNAME) $(TARSRC)
```

6 MapReduceFramework.cpp

```
1  #include <pthread.h>
2  #include <atomic>
3  #include <iostream>
4  #include <utility>
5  #include <algorithm>
6  #include <semaphore.h>
7  #include "MapReduceFramework.h"
8  #include "Barrier.h"
9
10 using std::cout;
11 using std::endl;
12 using std::vector;
13 using std::pair;
14
15 struct JobContext;
16 struct ThreadContext;
17
18 // Comparator.
19 bool comparePtrToPair(IntermediatePair a, IntermediatePair b)
20 { return a.first->operator<(*b.first); }
21
22 /**
23  * @brief Context of a thread.
24  */
25 typedef struct ThreadContext
26 {
27     // Job that spawned the thread.
28     JobContext *jobContext;
29     // Number of thread given upon initialization.
30     int threadNum;
31     // Input vector given by client.
32     InputVec inputVec;
33     // Intermediate vector.
34     vector<IntermediatePair> *interVec;
35     // Semaphore.
36     sem_t *sem;
37     // Atomic counter.
38     std::atomic<int> *atomicCounter;
39     // Output vector.
40     OutputVec *outputVec;
41     // Client.
42     const MapReduceClient *client;
43     // Barrier.
44     Barrier *barrier;
45     // Mutex.
46     pthread_mutex_t *mutex;
47
48     // Ctor.
49     ThreadContext(JobContext *_jobContext, int _threadNum, const InputVec &_inputVec, sem_t *_sem,
50                   std::atomic<int> *_atomicCounter, OutputVec &_outputVec, const MapReduceClient *_client,
51                   Barrier *_barrier, pthread_mutex_t *_mutex) :
52         jobContext(_jobContext), threadNum(_threadNum), inputVec(_inputVec),
53         interVec(new IntermediateVec()), sem(_sem), atomicCounter(_atomicCounter),
54         outputVec(&_outputVec), client(_client), barrier(_barrier), mutex(_mutex)
55     {}
56
57 } ThreadContext;
58
59 /**
```

```

60  * @brief Context of a job.
61  */
62  typedef struct JobContext
63  {
64      // Vector of threads alive within this job.
65      vector<ThreadContext *> *threads{};
66      pthread_t *threadArr;
67      // State of the current job.
68      JobState *state;
69
70      // Ctor for a JobContext instance. Receives _threads as pointer.
71      JobContext(vector<ThreadContext *> *_threads, pthread_t *_threadArr) : threads(_threads), threadArr(_threadArr)
72      {
73          // Inits state.
74          state = new JobState();
75          // Sets state.
76          state->stage = UNDEFINED_STAGE;
77          state->percentage = 0;
78      }
79
80      ~JobContext()
81      {
82          if (!threads->empty())
83          {
84              delete threads->back()->barrier;
85              delete threads->back()->mutex;
86              delete threads->back()->atomicCounter;
87              delete threads->back()->sem;
88          }
89          for (ThreadContext *tc:*threads)
90          {
91              delete tc->interVec;
92              delete tc;
93          }
94          delete threads;
95          delete[] threadArr;
96          delete state;
97      }
98  } JobContext;
99
100
101  /**
102   * @brief Emits pairs into context = intermediate vector.
103   */
104  void emit2(K2 *key, V2 *value, void *context)
105  {
106      auto *interVec = (vector<IntermediatePair> *) context;
107      auto p = IntermediatePair(key, value);
108      interVec->push_back(p);
109  }
110
111  void emit3(K3 *key, V3 *value, void *context)
112  {
113      auto curr_context = (ThreadContext *) context;
114      auto p = OutputPair(key, value);
115      curr_context->outputVec->push_back(p);
116  }
117
118  // Mapping
119  void mapPhase(ThreadContext *context)
120  {
121      context->jobContext->state->stage = MAP_STAGE;
122      vector<IntermediatePair> *interVec = context->interVec;
123      int oldValue;
124      // Use atomic to avoid race conditions.
125      while ((unsigned long int) context->atomicCounter->load() < context->inputVec.size())
126      {

```



```

128         oldValue = (*(context->atomicCounter))++;
129         InputPair currPair = context->inputVec.at(static_cast<unsigned int>(oldValue));
130         // Map each pair.
131         context->client->map(currPair.first, currPair.second, interVec);
132         // Update percentage.
133         context->jobContext->state->percentage = ((oldValue + 1) / (float) context->inputVec.size() * 100);
134     }
135 }
136
137 // Sorting.
138 void sortPhase(ThreadContext *context)
139 {
140     std::sort(context->interVec->begin(), context->interVec->end(), comparePtrToPair);
141 }
142
143 void shufflePhase(vector<IntermediateVec> *reduceQueue, ThreadContext *context)
144 {
145     context->jobContext->state->stage = REDUCE_STAGE;
146
147     // Throw all pairs into a single vector.
148     IntermediateVec newInter;
149     for (ThreadContext *tc: *context->jobContext->threads)
150     {
151         newInter.insert(newInter.begin(), tc->interVec->begin(), tc->interVec->end());
152     }
153     std::sort(newInter.begin(), newInter.end(), comparePtrToPair);
154
155     // Go over all pairs.
156     K2 *k2max = newInter.back().first;
157     IntermediatePair *currPair = &newInter.back();
158     IntermediateVec currVec;
159     while (!newInter.empty())
160     {
161         if (*currPair->first < *k2max)
162         {
163             if (pthread_mutex_lock(context->mutex) != 0)
164             {
165                 fprintf(stderr, "Shuffle: error on pthread_mutex_lock");
166                 exit(1);
167             }
168             reduceQueue->push_back(currVec);
169             sem_post(context->sem);
170             if (pthread_mutex_unlock(context->mutex) != 0)
171             {
172                 fprintf(stderr, "Shuffle: error on pthread_mutex_unlock");
173                 exit(1);
174             }
175             currVec = IntermediateVec();
176             k2max = currPair->first;
177         }
178         currVec.push_back(*currPair);
179         newInter.pop_back();
180         currPair = &newInter.back();
181     }
182
183     if (pthread_mutex_lock(context->mutex) != 0)
184     {
185         fprintf(stderr, "Shuffle: error on pthread_mutex_lock");
186         exit(1);
187     }
188     if (!currVec.empty())
189     {
190         reduceQueue->push_back(currVec);
191         sem_post(context->sem);
192     }
193     if (pthread_mutex_unlock(context->mutex) != 0)

```

```

196     {
197         fprintf(stderr, "Shuffle: error on pthread_mutex_unlock");
198         exit(1);
199     }
200
201 }
202
203 // Reducing
204 void reducePhase(vector<IntermediateVec> *reduceQueue, int reduceSize, ThreadContext *context)
205 {
206     while (!reduceQueue->empty())
207     {
208         sem_wait(context->sem);
209
210         if (pthread_mutex_lock(context->mutex) != 0)
211         {
212             fprintf(stderr, "Reduce: error on pthread_mutex_lock");
213             exit(1);
214         }
215
216         if (!reduceQueue->empty())
217         {
218             context->client->reduce(&reduceQueue->back(), context);
219             reduceQueue->pop_back();
220             context->jobContext->state->percentage = (reduceSize - reduceQueue->size()) / (float) reduceSize * 100;
221         }
222         else
223         {
224             sem_post(context->sem);
225         }
226
227         if (pthread_mutex_unlock(context->mutex) != 0)
228         {
229             fprintf(stderr, "Reduce: error on pthread_mutex_lock");
230             exit(1);
231         }
232     }
233 }
234
235 /**
236  * @brief The main function of each thread.
237  */
238 void threadMapReduce(ThreadContext *context)
239 {
240     mapPhase(context);
241     sortPhase(context);
242     context->barrier->barrier(); // waiting for unlock.
243     auto *reduceQueue = new vector<IntermediateVec>();
244     if (context->threadNum == 0)
245     {
246         shufflePhase(reduceQueue, context);
247     }
248     reducePhase(reduceQueue, reduceQueue->size(), context);
249     delete reduceQueue;
250 }
251
252 JobHandle startMapReduceJob(const MapReduceClient &client, const InputVec &inputVec, OutputVec &outputVec,
253                             int multiThreadLevel)
254 {
255     auto *sem = new sem_t();
256
257     // atomic counter to be used as input vec index.
258     auto *atomic_counter = new std::atomic<int>(0);
259     auto *threads = new vector<ThreadContext *>();
260     auto *threadArr = new pthread_t[multiThreadLevel];
261     auto *jobContext = new JobContext(threads, threadArr);
262     auto *barrier = new Barrier(multiThreadLevel);
263     auto *mutex = new pthread_mutex_t(PTHREAD_MUTEX_INITIALIZER);

```

```

264     for (int i = 0; i < multiThreadLevel; ++i)
265     {
266         ThreadContext *context = new ThreadContext(jobContext, i, inputVec, sem,
267             atomic_counter, outputVec, &client,
268             barrier, mutex);
269         threads->push_back(context);
270         pthread_create(threadArr + i, nullptr, (void (*)(void *)) threadMapReduce, context);
271     }
272     return jobContext;
273 }
274
275 void waitForJob(JobHandle job)
276 {
277     auto *jobContext = (JobContext *) job;
278     while (jobContext->state->stage != REDUCE_STAGE || jobContext->state->percentage < 100)
279     {
280     }
281 }
282
283 void getJobState(JobHandle job, JobState *state)
284 {
285     auto *jc = (JobContext *) job;
286     *state = *jc->state;
287 }
288
289 void closeJobHandle(JobHandle job)
290 {
291     waitForJob(job);
292     auto jobContext = (JobContext *) job;
293     for (int i = 0; i < (int)jobContext->threads->size(); ++i)
294     {
295         pthread_join(jobContext->threadArr[i], nullptr);
296     }
297     delete jobContext;
298 }

```